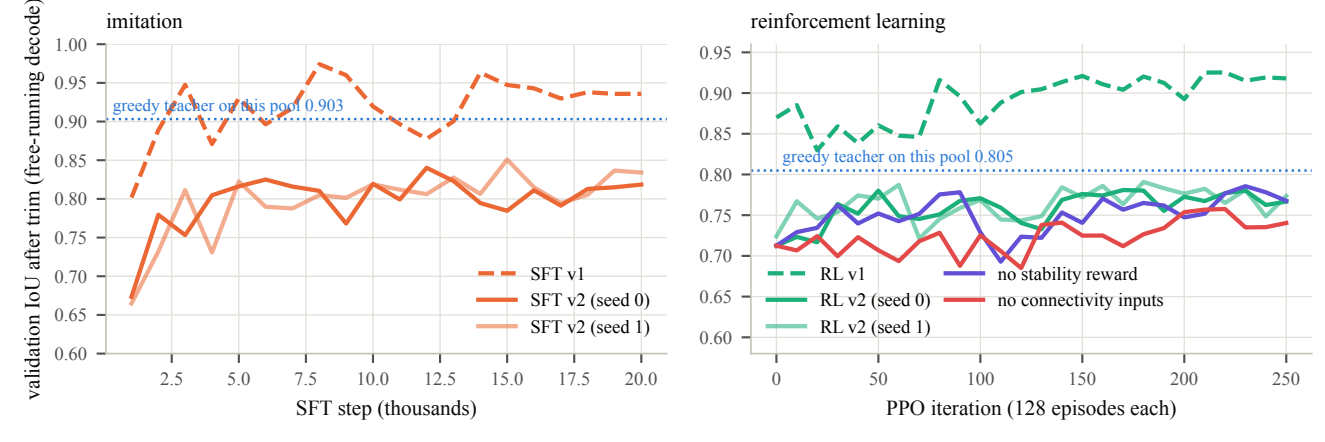
imitation
validation IoU after trim (free-running decode)
1.00
0.95
0.90
0.85
0.80
0.75
0.70
0.65
0.60
greedy teacher on this pool 0.903
SFT v1
SFT v2 (seed 0)
SFT v2 (seed 1)
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
SFT step (thousands)
reinforcement learning
0.95
0.90
0.85
0.80
0.75
0.70
0.65
0.60
greedy teacher on this pool 0.805
RL v1
RL v2 (seed 0)
RL v2 (seed 1)
no stability reward
no connectivity inputs
0
50
100
150
200
250
PPO iteration (128 episodes each)